

From project to product: competing with Muse and Instinct

October 2026

All of Muse and Instinct, running on your own computer. Full control. Real privacy. Zero token cost.

Shorter variants, for different surfaces:

- App Store / site hero: *"Your AI agent. Your computer. Your rules."*
- Comparison ads: *"Everything Muse and Instinct do, without sending them your life."*
- Pricing page: *"No tokens. No meter. No cloud required."*

1. Where we stand

LocalAIAgent 0.15.0 already matches most of Muse's and Instinct's feature list on a single computer (see `OVERVIEW.md` §5). It also has advantages neither offers:

- visible decision probabilities;
- an egress-free default;
- end-to-end encrypted agent-to-agent without a server;
- no running cost.

Four gaps remain between a strong project and a product people choose over them:

Gap	Why it matters to users	Muse/Instinct today
1. Always there	"It didn't reply because my laptop was asleep" kills the habit.	Cloud VMs never sleep.
2. Smart enough	Long tasks with a 4B model fail more often than with a frontier model.	Frontier models.
3. Reaches everything	1,500+ connectors and shopping partners vs our ~15 integrations.	Partner ecosystems, one-time cards.
4. Feels like a product	Terminal install, Python, permissions granted to "Terminal".	Polished apps on every device.

A fifth item is our moat, not a gap: **trust**. It has to be turned from architecture into something users can see (`TRUST_UX.md`) and something that is actually secure (`SECURITY_REVIEW.md` P0/P1).

2. Capabilities to build

Pillar A: Trust you can verify (our moat; do first)

Capability	What it is	Why it wins
Security P0/P1 fixes	Authenticated local API, data-flow taint, executable blocking, payload-visible approvals, per-target grants	Prerequisite for any public release.
Privacy Control Center	One screen: every capability, OS permission, last use, on/off, "turn off and forget"; trust presets (Observer / Assistant / Agent)	Users can see and <i>change</i> what the agent can do.
Egress ledger + weekly privacy receipt	Every byte that left the machine, with the exact content	Proves the local claim. Neither competitor can show this.
"Why?" on everything	Explanation card for every action and nudge: decision probabilities, sources, the rule that allowed it	Explainability as a feature, not a doc.
Agent firewall mode	A switch that blocks all outbound network except Ollama and explicitly approved destinations; enforced in-process and optionally with the OS firewall	"Airplane mode for your AI".
Open core + reproducible builds + third-party audit	Publish the agent core, sign releases, commission an audit	Trust from people who can check.
Hard limits	No binding agreements, no money movement, nothing irreversible by phone, ever	The opposite of Instinct's terms.

Pillar B: Always available, on every device (closes Gap 1)

Capability	What it is	Notes
Home node mode	Run the agent on an always-on machine (Mac mini, home PC, NAS, a small Linux box) and use laptops and phones as clients	The same package; <code>localagent serve --home</code> . Most households have or can buy a \$500 always-on box: still zero token cost.
Companion apps (iOS, Android)	A thin client: chat, approvals with full payload, voice, push notifications, share sheet, App Intents / Siri	Connect over Tailscale or our blind relay (below).
Blind relay	An optional rendezvous server that forwards end-to-end encrypted packets between your devices and your friends' agents. It can't read content and stores nothing beyond a short offline queue	It solves NAT and sleep for phones and peers without breaking privacy. The relay can be self-hosted.
Multi-device sync	Laptop ↔ home node sync of memory and settings (CRDT, E2E encrypted)	Start on the laptop, continue on the phone, as Muse does.
Wake-on-demand	Wake-on-LAN or <code>pmset</code> scheduling for proactive jobs; a "keep awake while charging" option	Mitigates sleep for laptop-only users.
Phone context (opt-in)	Location, focus mode and calendar from the phone app	Brings back Instinct's "the usual place" from the device that actually knows.

Pillar C: Smarter on small models (closes Gap 2)

Capability	What it is	Expected effect
Hardware-aware model tiers	Auto-select: 8 GB → 1.7B/4B; 16 GB → 4B/8B; 32 GB+ → 14B–32B (MLX/llama.cpp); a home node with a GPU runs larger models	Better quality for users with better hardware.

Capability	What it is	Expected effect
Fine-tuned tool-calling model	LoRA on Qwen-class 4–8B using our tool schemas and synthetic and corrected traces (from opted-in users only, locally)	Fewer wrong tool calls and malformed arguments.
Plan → act → verify loop	Explicit plans with checkpoints; verify each step's result before the next; retry with a narrower prompt	Long tasks stop drifting.
Learned recipes ("do it like last time")	A successful multi-step run (e.g. a Safeway price check) becomes a deterministic recipe the agent replays, with the model handling only the variable parts	Turns 20 fragile steps into one reliable one. Shareable as skills.
Real System-1 decision model	Distil the prototype + judge decisions and user corrections into a small Jev-style model (ModernBERT class)	Calibrated probabilities, ~10 ms, multilingual.
Parallel sub-agents	llama.cpp/Ollama parallel slots and a task queue with priorities; background tasks with progress	Muse-style "work on these three things".
Optional cloud, still in control	Any provider or the user's own key, with redaction (names, numbers, emails replaced before sending), per-request approval or a monthly budget	A frontier model when the user wants one, without giving up control.

Pillar D: Reaches everything (closes Gap 3)

Capability	What it is	Notes
MCP client	Plug in any Model Context Protocol server; each tool gets a tier (default <i>write</i>), untrusted-output marking, and appears in the Control Center	Thousands of existing connectors on day one, under our policy engine.
Signed skill marketplace	Skills with a permission manifest (folders, domains, tier), signatures, reviews; run sandboxed (and on Windows, using AppContainer)	Our answer to the "1,500 connectors in a week" story, with safety built in.
First-party connectors	Google Calendar/Drive, Outlook/Microsoft 365, Slack, Notion, Todoist, Home Assistant/HomeKit, Spotify, Apple Health export, bank/brokerage CSV import	Cover the top 90% of daily use.
Commerce with the user's own card controls	Virtual cards the user creates (Privacy.com, card-issuer single-use numbers, Apple Pay handoff) with a per-task cap; card details live in the vault and are typed by the tool layer, never seen by the model	Muse/Instinct-style checkout without a partner deal; payment remains a <i>danger</i> approval.
Merchant recipes	Maintained recipes for the top retailers, groceries, restaurants and travel sites; order tracking from email	Reliability where users shop most.
Calls on your behalf (opt-in, later)	Outbound calls with an explicit AI disclosure, a script approved before dialling, and a recording kept locally	Only after the legal review; never negotiating or agreeing to terms.

Pillar E: Feels like a product (closes Gap 4)

Capability	What it is
Native app shell	A signed and notarised macOS app with a menu-bar icon (Swift or Tauri), a Windows MSIX, Linux Flatpak. It bundles Python and talks to Ollama. Its own TCC identity fixes the "Terminal has Full Disk Access" problem.
One-click install, signed auto-update	Ollama and models fetched with progress, verified by digest. Delta updates. Rollback.
Real-time voice	Streaming STT + neural TTS (Kokoro/Piper) with barge-in, about 500 ms turn latency on Apple Silicon; voice notes over iMessage.

Capability	What it is
Avatar 2.0	Expressive 2D/3D avatar with visemes from the TTS phonemes; optional photoreal on hardware that can afford it.
Family and teams	Separate agents per person on one home node; shared calendars; parental controls.
Accessibility and languages	Multilingual UI and voice; screen-reader-first UI.

3. Phased plan

Phase	Theme	Main deliverables
1. Trustworthy beta (6–8 weeks)	Pillar A	P0+P1 security fixes; Privacy Control Center; egress ledger; payload approvals; off-by-default onboarding; native macOS shell with its own TCC identity; signed installer
2. Always there (8–10 weeks)	Pillar B	Home node mode; iOS companion app (chat, approvals, push); blind relay (open source, self-hostable); E2E sync
3. Smarter (8–12 weeks, overlapping)	Pillar C	Hardware tiers; plan-act-verify; learned recipes; fine-tuned tool model; parallel tasks; redacting cloud option
4. Everything (ongoing)	Pillar D	MCP client; skill marketplace with signed manifests; top-10 connectors; virtual-card checkout
5. Polish (ongoing)	Pillar E	Real-time voice; Android; Windows/Linux native shells; avatar 2.0; family accounts

Release gates per phase:

- the security review re-run with no open P0/P1;
- egress ledger tests ("this feature sends nothing" asserted in CI);
- an injection red-team suite passing;
- crash-free sessions above 99.5%.

4. Business model that keeps "zero token cost" honest

Tier	Price	What
Free	\$0	The full local agent, all core features, unlimited use: no tokens, no meter
Pro	One-time license or ~\$5/month	Companion apps and push via the blind relay, multi-device sync, premium voices, priority skill updates
Family	~\$10/month	Up to 6 people on one home node, shared planning
Home node	Hardware partner bundle	Pre-configured mini PC with the agent and models installed
Marketplace	Revenue share	Paid skills and merchant recipes from third parties

What we never sell: user data, model access to user data, or ads. Cloud model use, if any, is always on the user's own key.

5. Positioning against each competitor

Message	vs Meta Muse	vs Instinct
Privacy	"Muse needs your life on Meta's servers. We never see it."	"Instinct kept data after people disconnected. Ours is a folder you can delete."
Control	"Approvals with the full payload, and a ledger of everything that left."	"Never signs anything for you. No binding agreements, ever."
Cost	"No subscription required. Your computer does the work."	Same
Transparency	"See why it did everything, with the numbers."	"See the email that tried to hijack it, and that it said no."
Openness	"Open core, auditable, extensible with MCP and signed skills."	Same

6. Risks and how to handle them

Risk	Mitigation
Local models stay noticeably weaker on long tasks	Recipes and plan-verify for reliability; hardware tiers; optional cloud with redaction; focus marketing on daily tasks where small models do well.
The always-on story depends on a second device	Home node bundles; wake-on-demand; be honest that a laptop-only setup pauses when it sleeps.
Platform changes (macOS privacy, Messages DB format)	Native shell with proper entitlements; integration tests per macOS beta; prefer official APIs (EventKit, App Intents).
Merchant blocking of agents	Recipes that act as the user in their own browser; hand off to the user for checkout.
Security incident	P0/P1 first; red-team suite in CI; bug bounty; fast signed updates.